◆ PREVIOUS NEXT ▶

## 14.6 Ancillary Data

Ancillary data can be sent and received using the msg\_control and msg\_controllen members of the msghdr structure with the sendmsg and recvmsg functions. Another term for ancillary data is *control information*. In this section, we will describe the concept and show the structure and macros used to build and process ancillary data, but we will save the code examples for later chapters that describe the actual uses of ancillary data.

Figure 14.11 is a summary of the various uses of ancillary data we cover in this text.

Figure 14.11. Summary of uses for ancillary data.

Protocol	cmsg_level	cmsg_type	Description
IPv4	IPPROTO_IP	IP_RECVDSTADDR IP_RECVIF	Receive destination address with UDP datagram Receive interface index with UDP datagram
IPv6	IPPROTO_IPV6	IPV6_DSTOPTS IPV6_HOPLIMIT IPV6_HOPOPTS IPV6_NEXTHOP IPV6_PKTINFO IPV6_RTHDR IPV6_TCLASS	Specify destination options Specify hop limit Specify hop-by-hop options Specify next-hop address Specify packet information Specify routing header Specify traffic class
Unix domain	SOL_SOCKET	SCM_RIGHTS SCM_CREDS	Send/receive descriptors Send/receive user credentials

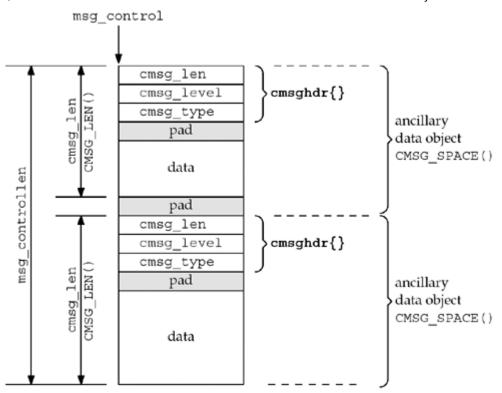
The OSI protocol suite also uses ancillary data for various purposes we do not discuss in this text.

Ancillary data consists of one or more *ancillary data objects*, each one beginning with a cmsghdr structure, defined by including <sys/socket.h>.

We have already seen this structure in <u>Figure 14.9</u>, when it was used with the <u>IP\_RECVDSTADDR</u> socket option to return the destination IP address of a received UDP datagram. The ancillary data pointed to by <u>msg\_control</u> must be suitably aligned for a <u>cmsghdr</u> structure. We will show one way to do this in <u>Figure 15.11</u>.

Figure 14.12 shows an example of two ancillary data objects in the control buffer. msg\_control points to the first ancillary data object, and the total length of the ancillary data is specified by msg\_controllen. Each object is preceded by a cmsghdr structure that describes the object. There can be padding between the cmsg\_type member and the actual data, and there can also be padding at the end of the data, before the next ancillary data object. The five CMSG\_xxx macros we describe shortly account for this possible padding.

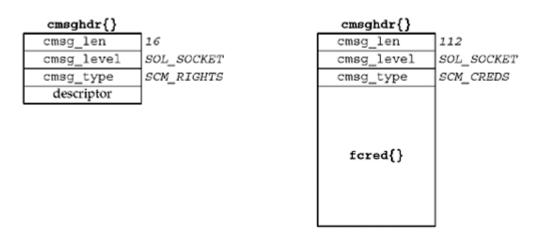
Figure 14.12. Ancillary data containing two ancillary data objects.



Not all implementations support multiple ancillary data objects in the control buffer.

<u>Figure 14.13</u> shows the format of the <u>cmsghdr</u> structure when used with a Unix domain socket for descriptor passing (<u>Section 15.7</u>) or credential passing (<u>Section 15.8</u>).

Figure 14.13. cmsghdr structure when used with Unix domain sockets.



In this figure, we assume each of the three members of the cmsghdr structure occupies four bytes and there is no padding between the cmsghdr structure and the actual data. When descriptors are passed, the contents of the cmsg\_data array are the actual descriptor values. In this figure, we show only one descriptor being passed, but in general, more than one can be passed (in which case, the cmsg\_len value will be 12 plus 4 times the number of descriptors, assuming each descriptor occupies 4 bytes).

Since the ancillary data returned by recvmsg can contain any number of ancillary data objects, and to hide the possible padding from the application, the following five macros are defined by including the <sys/socket.h> header to simplify the processing of the ancillary data:

```
#include <sys/socket.h>
#include <sys/param.h> /* for ALIGN macro on many implementations */
struct cmsghdr *CMSG_FIRSTHDR(struct msghdr *mhdrptr);
                                               Returns: pointer to first cmsghdr structure or NULL if no ancillary data
struct cmsghdr *CMSG_NXTHDR(struct msghdr *mhdrptr, struct cmsghdr *cmsgptr);
                                  Returns: pointer to next cmsghdr structure or NULL if no more ancillary data objects
unsigned char *CMSG_DATA(struct cmsghdr *cmsgptr);
                                              Returns: pointer to first byte of data associated with cmsghdr structure
unsigned int CMSG_LEN(unsigned int length);
                                                     Returns: value to store in cmsg_len given the amount of data
unsigned int CMSG_SPACE(unsigned int length);
                                               Returns: total size of an ancillary data object given the amount of data
```

POSIX defines the first three macros; RFC 3542 [Stevens et al. 2003] defines the last two.

These macros would be used in the following pseudocode:

CMSG\_FIRSTHDR returns a pointer to the first ancillary data object, or a null pointer if there is no ancillary data in the msghdr structure (either msg\_control is a null pointer or cmsg\_len is less than the size of a cmsghdr structure). CMSG\_NXTHDR

returns a null pointer when there is not another ancillary data object in the control buffer.

Many existing implementations of CMSG\_FIRSTHDR never look at msg\_controllen and just return the value of cmsg\_control. In Figure 22.2, we will test the value of msg\_controllen before calling this macro.

The difference between CMSG\_LEN and CMSG\_SPACE is that the former does not account for any padding following the data portion of the ancillary data object and is therefore the value to store in cmsg\_len, while the latter accounts for the padding at the end and is therefore the value to use if dynamically allocating space for the ancillary data object.

◆ PREVIOUS NEXT ▶