

Dave Pacheco's Blog

Systems software engineering

About me

Performance Puzzler: The Slow Server

July 8, 2019 by [dap](#)

In the spirit of the [TCP Puzzlers](#), I want to present a small distributed systems performance puzzler: **In a distributed system with a fixed-number of servers and a fixed concurrency, what happens when one of the backend servers gets slow?** First, let's pick apart what this means.

Suppose you have:

- a service with a fixed number of backend servers — say, 100 servers — each exactly equivalent. You could imagine this as a load balancing service having 100 backends or a key-value store having 100 well-distributed shards.
- a fixed number of clients — say, 10,000. Each client makes one request to the system at a time. When each request completes, the client makes another one. The 10,000 clients could be 1,000 clients each having a 10-way thread pool or 10 event-oriented clients each capable of 1,000 concurrent requests — either way, the total concurrency is 10,000. It doesn't matter of the clients and server communicate via RPC, HTTP, or something else.

with the following assumptions (which we'll revisit later):

- All requests take about the same amount of work to process, both on the client-side and server-side.
- The per-request cost outside the server is negligible. (We're going to ignore client-to-server transit time as well as client processing time.)
- The backend servers are infinitely scalable. That is, we'll assume that if they take 1ms to process 1 request, they'll also take 1ms to process 10 requests issued at the same time.
- Each request is dispatched to one of the servers uniformly at random.

To start with, suppose each server takes 1ms to complete each request. What's the throughput of the entire system? What will be the per-server throughput?

There are a few ways to analyze this:

- From the client perspective, we know there will be 10,000 requests outstanding to the system at any given instant. (Remember, there are 10,000 clients, each making a single request at a time, and we're assuming zero per-request costs on the client side and in transit.) After 1ms, those requests have all completed and another 10,000 requests are running. At this rate, the system will complete (10,000 requests per millisecond) times (1,000 milliseconds per second) = 10 million requests per second. Intuitively, with all servers equal, we'd expect each server to be handling 1% of these, or 100,000 requests per second.
- From the server perspective: with 10,000 outstanding requests at all times and 100 servers, we'd expect about 100 requests outstanding per server at all times. Each takes 1ms, so there would be 100,000 completed per server per second (which matches what we said above).

Now, **what happens when one of the servers becomes slow?**

Specifically, suppose all of a sudden one of the servers starts taking 1s to process every request. What happens to overall throughput? What about per-server throughput? (Make a guess before reading on!)

The analysis is a bit more complex than before. From the client perspective, there will still always be 10,000 requests outstanding. 99% of requests will land on a fast server and take 1ms, while the remaining 1% will land on the slow server and take 1s. The expected latency for any request is thus $0.99 * 1\text{ms} + 0.01 * 1000\text{ms} = 10.99\text{ms}$. If each of the 10,000 clients completes (1 request / 10.99ms) times (1,000 milliseconds per second), we get an expected throughput of about 91 requests per client, or 909,918 requests per second. (That's a degradation of 91% from the original 10 million requests per second!)

Above, we assumed that each server had the same throughput, but that's not so obvious now that one of them is so slow. Let's think about this another way on the client side: what's the probability at any given time that a client has a request outstanding to the slow server? I find it easiest to imagine the assignment being round-robin instead of random.

Suppose clients issued requests to each of the 99 fast servers, then the slow server, and then started again with the first fast server. In that case, the first 99 requests would take 99ms, and the last request would take 1s. The client would have a request outstanding to the slow server for $1,000\text{ms} / 1,099\text{ms}$ or about 91% of the time. There's a more rigorous explanation below, but I think it's intuitive that the random distribution of work would behave the same way on average.

In that case, we could also expect that 91% of clients (or 9100 clients) have a request outstanding to the slow server at any given time. Each of these takes 1 second. The result is a throughput of 9,100 requests per second on the slow server.

What about the fast server? We'd expect that each fast server has 1/99 of the remaining requests (9% of requests) at any given time, which works out to 9.1 requests per server. At 1ms per request, these servers are doing (drum roll) 9,100 requests per second. The same as the slow server! Is that what you guessed? I didn't.

This isn't an accident of the specific numbers I picked. If you run the same analysis algebraically instead of with concrete numbers, you'll find that the general result holds: in this system, when one server becomes slow, the throughput of every server remains the same.

A (more) formal proof

Let's define some notation:

- C : the total client concurrency (10,000 in our example)
- N : the number of servers (100 in our example)
- S_i : server i (where i ranges from 1 to N)
- L_i : the latency (in seconds) of requests to server S_i . In our example, L_1 would be 1 second, while $L_j = 0.001$ for $j > 1$.
- C_i : the number of requests outstanding on server S_i at any given time.
- T_i : the throughput of server S_i (in requests per second).
- $Pr(c_i)$: the probability that a particular client has a request outstanding to server S_i at any given time.
- $Pr(R_i)$: the probability that a newly-issued request will be issued to server S_i .

We're looking for a formula for T_i . For a server that can handle one request a time, the throughput is simply the inverse of the latency ($1/L_i$). We said that our server was infinitely scalable, which means it can execute an arbitrary number of requests in parallel. Specifically, it would be executing C_i requests in parallel, so the throughput T_i is:

$$T_i = C_i \cdot \frac{1}{L_i}$$

Now, how do we calculate C_i , the concurrency of requests at each server? Well, since each client has exactly one request outstanding at a time, C_i is exactly the number of clients that have a request outstanding to server S_i . Since the clients operate independently, that's just:

$$C_i = C \cdot Pr(c_i)$$

To calculate c_i , we need the percentage of time that each client has a request outstanding to S_i . We define a probability space of events (t, i) indicating that at millisecond timestamp t , the client has a request

outstanding to server S_i . By simply counting the events, we can say that the probability that a client has a request outstanding to server S_i is:

$$Pr(c_i) = \frac{L_i \cdot Pr(R_i)}{\sum_j L_j \cdot Pr(R_j)}$$

Now, since it's equally likely that we'll assign any given request to any server:

$$Pr(R_i) = \frac{1}{N}$$

that means:

$$\begin{aligned} Pr(c_i) &= \frac{L_i \cdot \frac{1}{N}}{\sum_j L_j \cdot \frac{1}{N}} \\ &= \frac{L_i}{\sum_j L_j} \end{aligned}$$

All we're saying here is that the fraction of time each client spends with a request outstanding to a particular server is exactly that server's latency divided by the latency of all servers put together. This makes intuitive sense: if you have 20 servers that all take the same amount of time, each client would spend 5% (1/20) of its time on each server. If one of those servers takes twice as long as the others, it will spend 9.5% (2/21) of its time on that one (almost twice as much as before) and 4.8% (1/21) on the others.

With this, we can go back to C_i :

$$\begin{aligned} C_i &= C \cdot \frac{L_i}{\sum_j L_j} \\ C_i &= \frac{L_i}{\sum_j L_j} \end{aligned}$$

and finally back to T_i :

$$\begin{aligned}T_i &= C_i \cdot \frac{1}{L_i} \\&= \frac{L_i}{\sum_j L_j} \cdot \frac{1}{L_i} \\&= \frac{C}{\sum_j L_j}\end{aligned}$$

and we have our result: the throughput at each server does not depend at all on the latency of that server (or any other server, for that matter)!

We can plug in some numbers to sanity check. In our example, we started with:

$$\begin{aligned}C &= 10,000 \\ \sum_j L_j &= 1.099\end{aligned}$$

and we get the results:

$$T_i = 10,000/1.099 = 9,099rps$$

which matches what we said above.

Simulating the behavior

One might find this result hard to believe. I wrote a simple [simulator](#) to demonstrate it. The simulator maintains a virtual clock with a 1ms resolution. To start with, each client issues a request to a random server. Each virtual millisecond, the simulator determines which requests have completed and has the corresponding clients make another request. This runs for a fixed virtual time, after which we determine the total throughput and the per-server throughput.

The simulator as-is hardcodes the the scenario that I described above (1 server with 1s per request, 99 servers with 1ms per request) and a 10-minute simulation. For the specific numbers I gave earlier, the results are:

```

$ node sim.js
simulated time:          600000 ms
total client concurrency: 10000
servers:
    1 server that completes requests with latency 1000 ms
    99 servers that complete requests with latency 1 ms (

server_0    5455039 requests ( 9092 rps)
server_1    5462256 requests ( 9104 rps)
server_2    5459316 requests ( 9099 rps)
server_3    5463211 requests ( 9105 rps)
server_4    5463885 requests ( 9106 rps)
server_5    5456999 requests ( 9095 rps)
...
server_95    5457743 requests ( 9096 rps)
server_96    5459207 requests ( 9099 rps)
server_97    5458421 requests ( 9097 rps)
server_98    5458234 requests ( 9097 rps)
server_99    5456471 requests ( 9094 rps)
overall:    545829375 requests (909715 rps, expect 9097 rps)

```

In my simulations, `server_0` is generally a bit slower than the others, but within 0.2% of the overall average. The longer I run the simulation, the closer it gets to the mean. Given that the slow server is three orders of magnitude slower per request (1s vs. 1ms), I'd say it's fair to conclude that the per-server throughput does not vary among servers when one server is slow.

Conclusions

The main result here is that **in this system, the per-server throughput is the same for all servers, even when servers vary significantly in how fast they process requests.** While the implementations of servers may be independent (i.e., they share no common hardware or software components), their behavior is not independent: the performance of one server depends on that of other servers! The servers have essentially been coupled to each other via the clients.

This result implies another useful one: **in a system like this one, when overall throughput is degraded due to one or more poorly-performing servers, you cannot tell from throughput alone which server is slow — or even how many slow servers there are!** You could determine which server was slow by looking at per-server latency.

Note too that even when this system is degraded by the slow server, the vast majority of requests complete very quickly. At the same time, clients spend the vast majority of their time waiting for the slow server. (We've come to refer to this at Joyent as "getting Amdahl'd", after Amdahl's Law.)

If you found this result surprising (as I did), then another takeaway is that the emergent behavior of even fairly simple systems can be surprising. This is important to keep in mind when making changes to the system (either as part of development or in production, as during incident response). Our intuition often leads us astray, and there's no substitute for analysis.

What about those assumptions?

The assumptions we made earlier are critical to our result. Let's take a closer look.

- *"There are a fixed number of clients."* This is true of some systems, but certainly not all of them. I expect the essential result holds as long as client count is not a function of server performance, which probably is true for many systems. However, if client software responds to poor performance by either reducing concurrency (figuring they might be overloading the server) or increasing it (to try to maintain a given throughput level), the behavior may well be different.
- *"All requests take about the same amount of work to process, both on the client-side and server-side."* This is true of some systems, but not others. When the costs differ between requests, I'd expect the actual per-client and per-request throughput to vary and make the result shown here harder to see, but it doesn't change the underlying result.

- *“The per-request cost outside the server is negligible.”* This isn’t true of real systems, but makes the math and simulation much simpler. I expect the results would hold as long as the per-request client and transit costs were fixed and small relative to the server cost (which is quite often true).
- *“The backend servers are infinitely scalable. That is, we’ll assume that if they take 1ms to process 1 request, they’ll take 1ms to process 10 requests issued at the same time.”*. This is the most dubious of the assumptions, and it’s obviously not true in the limit. However, for compute-bound services, this is not an unreasonable approximation up to the concurrency limit of the service (e.g., the number of available threads or cores). For I/O bound services, this is also a reasonable approximation up to a certain point, since disks often work this way. (For disks, it may not take much longer to do several read or write I/Os — even to spinning disks — than to do one such I/O. The OS can issue them all to the disk, and the disk can schedule them so they happen in one rotation of the platter. It’s not always so perfect as that, but it remains true that you can increase the concurrent I/O for disks up to a point without significantly increasing latency.)
- *“Each request is dispatched to one of the servers uniformly at random.”* Many client libraries (e.g., for RPC) use either a random-assignment policy like this one or a round-robin policy, which would have a similar result. For systems that don’t, this result likely won’t hold. In particular, a more sophisticated policy that keeps track of per-server latency might prefer the faster servers. That would send less work to the slower server, resulting in better overall throughput and an imbalance of work across servers.

As you can see, the assumptions we made were intended to highlight this particular effect — namely, the way a single slow server becomes a bottleneck in clients that affects throughput on other servers. If you have a real-world system that looks at all similar to this one and doesn’t have a more sophisticated request assignment policy, then most likely this effect is present to some degree but it maybe harder to isolate relative to other factors.

[← Visualizing PostgreSQL Vacuum Progress](#) [Modifying USDT providers with translated arguments →](#)

4 thoughts on “Performance Puzzler: The Slow Server”

Patrick domack says:

July 10, 2019 at 10:30 am

That is a lot of math to prove it. Hope it helps someone. But the other missing part of the requirements or expectations as you explained later was the round Robin or random selection method on the lb.

I normally always set using latency or weighted least connections or least connections to avoid this issue. Have issues also with http using sticky sessions or keepalives. But if it's an internal backend those can be turned off at the issue of overhead. I believe haproxy handles it per transaction instead of connection so that can help also.

So many options to turn loadbalancers and most people will not use, but in cases they are very useful.

Now if VMware had a better option than just round Robin iops.

dap says:

July 10, 2019 at 11:17 am

> I normally always set using latency or weighted least connections or least connections to avoid this issue. Have issues also with http using sticky sessions or keepalives. But if it's an internal backend those can be turned off at the issue of overhead. I believe haproxy handles it per transaction instead of connection so that can help also.

> So many options to turn loadbalancers and most people will not use, but in cases they are very useful.

Thanks for mentioning that! Absolutely. The Google SRE book goes into some detail about different policies to address situations like these. I've run into this problem in a sharded, stateful service. In that case, we didn't really have a choice about where to send the traffic because the state needed by the request lives in the slow cluster.

Johannes Rudolph says:

July 12, 2019 at 1:54 am

Thanks for the interesting post.

Here's another explanation (that basically follows your math) for that syncing of throughput between slow and fast servers:

Let's just look at a single client out of those 10000 and assume round robin load balancing. Its workload for the first 100 requests is to make one request to every server sequentially. I.e. by assumption each one of these requests depends on the previous one. If we look at the whole process, this is now a long sequential process where the total throughput is dictated by the sum of all execution times (sum of latencies to all servers), which in this case is dominated by the slow server latency, which is the bottleneck of the process. This process is then executed in a loop and concurrently by 10000 clients, which doesn't change the result.

This result is a direct consequence of the measuring method which assumes a fixed number of ongoing concurrent requests (which can be realistic for some workloads / client implementations and certainly is and was a well-known e.g. for HTTP load testing tools) instead of a fixed rate of incoming requests. A better measuring logic would therefore be to sample requests and either make sure that requests are scheduled independently of other ongoing requests or correct for the fact that some requests were executed later than planned in the sampling schedule (<https://github.com/HdrHistogram/HdrHistogram#corrected-vs-raw-value-recording-calls> does that correcting).

Titus says:

July 12, 2019 at 5:09 am

Thanks for the post!

It seems to me that the real problem is the round robin load balancing assumption. You would get much better throughput in this system by balancing with a different method, for example, picking a server to which to send the request by taking 1/3 least busy servers and picking a random one from these.

> I've run into this problem in a sharded, stateful service. In that case, we didn't really have a choice about where to send the traffic because the state needed by the request lives in the slow cluster.

How did you approach / solve the issue in the end? My first thought is to set aggressive timeouts on the server side to notice the faulty server / cluster

quickly and fix it.

Comments are closed.

Categories

Cloud Analytics (9)

DTrace (11)

Fishworks (13)

Joyent (28)

Manta (3)

Node.js (21)

Personal (3)

SmartOS (17)

Solaris (4)

Uncategorized (5)

[Follow me on Twitter](#)

[\(@dapsays \)](#)

[Follow **davepacheco** on](#)

[Github](#)

GitHub

javascriptlint (highly
configurable linter)

jsstyle (JavaScript style
checker)

node-vasync (Observable
asynchronous control flow)

node-verror (Richer JavaScript
errors)

node-jsprim (Utilities for
primitive types)

node-extsprintf (extended
sprintf)

node-getopt (getopt(3C)-style
option parser)

node-tab (Unix-style tables for
command-line utilities)

node-strsplit (Split a string by a
regular expression)

node-stackvis (Profiling
visualization tools)

portsnoop (Trace event port
activity)

smfgen (Generate SMF
manifests from a JSON
description)

mod_usdt (DTrace provider for
Apache)

Recent Posts

[Shell redirection example](#)

[Modifying USDT providers with translated arguments](#)

[Performance Puzzler: The Slow Server](#)

[Visualizing PostgreSQL Vacuum Progress](#)

[TCP puzzlers](#)

[Programming language debuggability](#)

[Debugging enhancements in Node 0.12](#)

[Understanding DTrace ustack helpers](#)

[Tracing Node.js add-on latency](#)

[Stopping a broken program in its tracks](#)

[Node.js in production: runtime log snooping](#)

[Kartlytics: Applying Big Data Analytics to Mario Kart 64](#)

[Fault tolerance in Manta](#)

[Inside Manta: Distributing the Unix shell](#)

[Debugging dynamic library dependencies on illumos](#)

[illumos tools for observing processes](#)

[OSCON Slides](#)

[NodeConf slides](#)

[ACM Turing Centenary Celebration](#)

[Debugging Node.js in Production \(Fluent slides\)](#)

Debugging RangeError from a core dump

Profiling Node.js

Managing Node.js dependencies with shrinkwrap

Playing with Node/V8 postmortem debugging

Where does your Node program spend its time?

USDT Providers Redux

Node.js/V8 postmortem debugging

Surge 2011

New metrics on no.de

JavaScript Lint on SmartOS

Distributed Web Architectures @ SF Node.js Meetup

OSCON Slides

Heatmap coloring

Heatmaps and more heatmaps

Presenting at OSCON Data 2011

Example: HTTP request latency and garbage collection

Welcome to Cloud Analytics

Tonight at 6: Solving Big Problems (with Cloud Analytics)

Joining Joyent

Leaving Oracle

SS7000 Software Updates

Replication for disaster recovery

Another detour: short-circuiting cat(1)

A ZFS Home Server

Replication in 2010.Q1

Remote Replication

Introduction

Threshold alerts

Anatomy of a DTrace USDT
provider

2009.Q2 Released

Compression followup

Archive

September 2019 (1)

July 2019 (2)

May 2019 (1)

August 2016 (1)

February 2015 (2)

November 2013 (1)

October 2013 (3)

August 2013 (1)

July 2013 (1)

June 2013 (1)

September 2012 (1)

August 2012 (1)

July 2012 (2)

June 2012 (1)

May 2012 (2)

April 2012 (1)

February 2012 (1)

January 2012 (2)

December 2011 (1)

October 2011 (2)

August 2011 (3)

July 2011 (1)

June 2011 (1)

May 2011 (2)

March 2011 (2)

January 2011 (1)

November 2010 (1)

September 2010 (4)

May 2010 (1)

April 2010 (1)

March 2010 (1)

August 2009 (1)

May 2009 (1)

April 2009 (1)

March 2009 (2)

February 2009 (1)

November 2008 (2)

February 2008 (1)

July 2006 (1)

June 2006 (3)

Meta

[Log in](#)

[Entries RSS](#)

[Comments RSS](#)

[WordPress.org](#)

2020 Dave Pacheco's Blog | Powered by WordPress | Theme Mon Cahier by Bluelime Media
